

## Project 3: Collaboration and Competition

In this report, I am going to present about the environment and the algorithms that I have used to solve collaboration and competition problem where the agents must bounce ball back and forth while not dropping or sending ball out of bounds.

### Environment

We work with Unity ML Environment, [Tennis](#), in this project.

In this environment, there are two agents which control rackets to bounce a ball over a net. The agent receives +0.1 reward if it manages to hit ball over the net without dropping or hitting out of bounds. For dropping or hitting out of bounds, it receives -0.01 .

The observation space consists of 24 variables representing position and velocity of the ball and racket. Each action is a vector with two numbers, corresponding to movement towards or away from the net, and jumping. The action vector should be a number between 1 and -1.

The environment is deemed solved if the agents get an average score of +0.5 over 100 consecutive episodes.

Given below are the characteristics of the agents and the environment.

- Unity brain name: TennisBrain
- Number of Visual Observations (per agent): 0
- Vector Observation space type: continuous
- Vector Observation space size (per agent): 8
- Number of stacked Vector Observation: 3
- Vector Action space type: continuous
- Vector Action space size (per agent): 2

### Learning Algorithm

To solve this reinforcement learning problem, I am using a Deep Deterministic Policy Gradients (DDPG) with modification to make it suitable for multiagent environment.

#### Deep Deterministic Policy Gradient (DDPG)

**DDPG** is an actor-critic algorithm that extends **DQN** to work in continuous spaces. Here, we use two deep neural networks, one as actor and the other as critic. Similar network architectures are used for both actor and critic. **ADAM** optimizer is used with **learning rates 0.0001** and **0.001** for actor and critic, respectively. And the **discount factor** used is **0.99**.

```
GAMMA = 0.99          # discount factor
LR_ACTOR = 1e-4        # learning rate of the actor
LR_CRITIC = 1e-3       # learning rate of the critic
```

## Neural Network Architecture

### Actor

State → BatchNorm → 128 → ReLU → 64 → ReLU → BatchNorm → action → tanh

### Critic

State → BatchNorm → 128 → Relu → 64 → Relu → action

## Pytorch Implementation

### Actor

```
self.fc1 = nn.Linear(state_size, 128)
self.bn1 = nn.BatchNorm1d(128)
self.fc2 = nn.Linear(128, 64)
self.bn2 = nn.BatchNorm1d(64)
self.fc3 = nn.Linear(64, 2)
self.bn3 = nn.BatchNorm1d(2)

...

x = self.fc1(state)
x = F.relu(x)
x = self.bn1(x)
x = self.fc2(x)
x = F.relu(x)
x = self.bn2(x)
x = self.fc3(x)
x = self.bn3(x)
x = F.tanh(x)
```

### Critic

```
self.bn0 = nn.BatchNorm1d(state_size)
self.fcs1 = nn.Linear(state_size, fcs1_units)
self.bn1 = nn.BatchNorm1d(fcs1_units)
self.fc2 = nn.Linear(fcs1_units+action_size, fc2_units)
self.bn2 = nn.BatchNorm1d(fc2_units)
self.fc3 = nn.Linear(fc2_units, 1)

...

x = self.bn0(state)
x = self.fcs1(x)
x = F.relu(x)
x = torch.cat((x, action), dim=1)
x = self.fc2(x)
x = F.relu(x)
x = self.fc3(x)
```

## Experience Replay

We store the last one million experience tuples (S,A,R,S') into a data container called **Replay Buffer** from which we sample a **mini batch of 1024** experiences. This batch ensures that the experiences are independent and stable enough to train the network.

```
BUFFER_SIZE = int(1e6) # replay buffer size
BATCH_SIZE = 1024      # minibatch size
```

## Soft Target Updates

In order to calculate the target values for both actor and critic networks, we use **Soft Target Update** strategy.

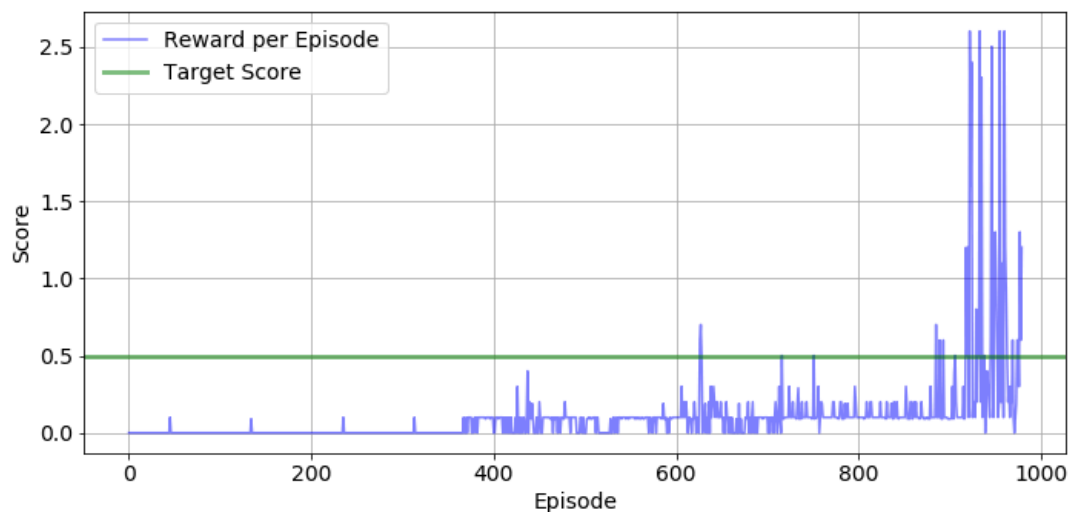
```
TAU = 1e-3
```

```
# for soft update of target parameters
```

In order for the agents to explore the entire environment, replay memory buffer is shared between the agents. Yet, `|| >` each agent has its own actor and critic networks to train.

## Plot of Rewards

After tuning the parameters and tweaking the network by changing the hidden layers size, I could solve the problem in **980 episodes**. The plot below shows the rewards per episode and the target.



## Result

Result can be seen in [tennis\\_result.gif](#)

## Ideas for Future Work

- After training, when I checked the performance, it seems that the agents are imitating each other. Reason could be the symmetry of the environment. Multi-Agent DDPG would be apt for environment like [Soccer](#). So, I am planning to implement and see the performance of MADDPG in Soccer environment.